

Google Cloud Skills Boost for Partners

[← Main menu](#)

05
Text Prompt
Engineering Techniques
Course · 8 hours ✓ Complete

[Course overview](#)

Text Prompt
Engineering
Techniques

- ✓ Get Started with Vertex AI Studio
- ✓ Getting Started with Google Generative AI Using the Gen AI SDK
- ✓ Common Generative AI Use Cases
- ✓ Prompt Design Strategies
- ✓ Generative AI with Vertex AI: Text Prompt Design: Challenge Lab

[Your Next Steps](#)[Introduction to Generative AI](#) > [Course](#) > [Text Prompt Engineering Techniques](#) >

Start Lab

01:00:00

—/100

Getting Started with Google Generative AI Using the Gen AI SDK

Lab 1 hour No cost Introductory

★★★★★ [Rate Lab](#)

This lab may incorporate AI tools to support your learning.

[Lab instructions and tasks](#)

GSP1209

Overview

Objectives

Setup and requirements

Task 1. Open the notebook in Vertex AI Workbench

Task 2. Set up the notebook

Task 3. Interact with the model

Task 4. Configure and control the model

Task 5. Manage the model interaction

[Task 6. Advanced features](#)

Congratulations!

GSP1209



Google Cloud Self-Paced Labs

Overview

The [Google Gen AI SDK](#) provides a unified interface to Google's generative AI API services. This SDK simplifies the process of integrating generative AI capabilities into applications and services, enabling developers to leverage Google's advanced AI models for various tasks. In this lab, you explore the Google Gen AI SDK, learning to connect to AI services, send diverse prompts, and fine-tune responses from Gemini. You also get hands-on experience with more advanced techniques to prepare you to leverage the power of generative AI for your own projects.

Objectives

In this lab, you learn how to use the Google Gen AI SDK for Python to interact with Google's generative AI services and models, including Gemini. You cover the following:

- Installing the Gen AI SDK.
- Connecting to an API service.
- Sending text and multimodal prompts.
- Setting system instructions.
- Configuring model parameters and safety filters.
- Managing model interactions (multi-turn chat, content streaming, asynchronous requests).
- Using advanced features (token counting, context caching, function calling, batch prediction, text embeddings).

Prerequisites

Before starting this lab, you should be familiar with:

- Basic Python programming.
- General API concepts.
- Running Python code in a Jupyter notebook on [Vertex AI Workbench](#).

Setup and requirements

Before you click the Start Lab button

Read these instructions. Labs are timed and you cannot pause them. The timer, which starts when you click **Start Lab**, shows how long Google Cloud resources are made available to you.

This hands-on lab lets you do the lab activities in a real cloud environment, not in a simulation or demo environment. It does so by giving you new, temporary credentials you use to sign in and access Google Cloud for the duration of the lab.

To complete this lab, you need:

- Access to a standard internet browser (Chrome browser recommended).

Note: Use an Incognito (recommended) or private browser window to run this lab. This prevents conflicts between your personal account and the student account, which may cause extra charges incurred to your personal account.

- Time to complete the lab—remember, once you start, you cannot pause a lab.

Note: Use only the student account for this lab. If you use a different Google Cloud account, you may incur charges to that account.

How to start your lab and sign in to the Google Cloud console

1. Click the **Start Lab** button. If you need to pay for the lab, a dialog opens for you to select your payment method. On the left is the Lab Details pane with the following:

- The Open Google Cloud console button
- Time remaining
- The temporary credentials that you must use for this lab
- Other information, if needed, to step through this lab

2. Click **Open Google Cloud console** (or right-click and select **Open Link in Incognito Window** if you are running the Chrome browser).

The lab spins up resources, and then opens another tab that shows the Sign in page.

Tip: Arrange the tabs in separate windows, side-by-side.

Note: If you see the **Choose an account** dialog, click **Use Another Account**.

3. If necessary, copy the **Username** below and paste it into the **Sign in** dialog.

"Username"



You can also find the Username in the Lab Details pane.

4. Click **Next**.

5. Copy the **Password** below and paste it into the **Welcome** dialog.

"Password"



You can also find the Password in the Lab Details pane.

6. Click **Next**.

Important: You must use the credentials the lab provides you. Do not use your Google Cloud account credentials.


Note: Using your own Google Cloud account for this lab may incur extra charges.

7. Click through the subsequent pages:

- Accept the terms and conditions.
- Do not add recovery options or two-factor authentication (because this is a temporary account).
- Do not sign up for free trials.

After a few moments, the Google Cloud console opens in this tab.

Note: To access Google Cloud products and services, click the **Navigation menu** or type the service or product name in the **Search** field.



Task 1. Open the notebook in Vertex AI Workbench

1. In the Google Cloud console, on the **Navigation menu** () , click **Vertex AI > Workbench**.
2. Find the **Workbench instance name** instance and click on the **Open JupyterLab** button.

The JupyterLab interface for your Workbench instance opens in a new browser tab.

Note: If you do not see notebooks in JupyterLab, please follow these additional steps to reset the instance:

1. Close the browser tab for JupyterLab, and return to the Workbench home page.
2. Select the checkbox next to the instance name, and click **Reset**.
3. After the **Open JupyterLab** button is enabled again, wait one minute, and then click **Open JupyterLab**.

Task 2. Set up the notebook

1. Open the notebook name file.
2. In the **Select Kernel** dialog, choose **Python 3** from the list of available kernels.
3. Run through the **Getting Started** section of the notebook. The Project ID and Location are pre-configured for you.

Note: If you experience a 429 response from any of the notebook cell executions, wait one minute before running the cell again to proceed.

Click **Check my progress** to verify the objective.



Import libraries and set up the notebook

Check my progress

Task 3. Interact with the model

For more information about all AI models and APIs on Vertex AI, refer to [Google Models](#) and [Model Garden](#).

Choose a model

- Run the **Choose a model** section of the notebook.

Send text prompts

Use the `generate_content` method to generate responses to your prompts. You can pass text to `generate_content`, and use the `.text` property to get the text content of the response.

- Run the **Send text prompts** section of the notebook.

Send multimodal prompts

You can include text, PDF documents, images, audio and video in your prompt requests and get text or code responses.

You can also pass the file URL in `Part.from_uri` in the request to the model directly.


- Run the **Send multimodal prompts** section of the notebook.

Set the system instructions

The [system instructions](#) allow you to control model behavior. Set the system instruction to give the model additional context to understand the task, provide more customized responses, and adhere to guidelines over the user interaction.

- Run the **Set system instruction** section of the notebook.

Click **Check my progress** to verify the objective.



Task 4. Configure and control the model

Configure model parameters

You can include parameter values in each call that you send to a model to control how the model generates a response. To learn more, refer to [experimenting with parameter values](#).

- Run the **Configure model parameters** section of the notebook.

Configure safety filters

Task 5. Manage the model interaction

Generate content stream

By default, the model returns a response after completing the entire generation process. You can also use the `generate_content_stream` method to stream the response as it is being generated. The model returns chunks of the response as they are generated.

- Run the **Generate content stream** section of the notebook.

Send asynchronous requests

You can send asynchronous requests using the `client.aio` module. This module exposes all the analogous async methods that are available on `client`.

For example, `client.aio.models.generate_content` is the async version of `client.models.generate_content`.

- Run the **Send asynchronous requests** section of the notebook.

Count tokens and compute tokens

You can use the `count_tokens` method to calculate the number of input tokens before sending a request to the Gemini API. Refer to the [List and count tokens](#) page for details.


Count tokens

- Run the **Count tokens** section of the notebook.

Compute tokens

- Run the **Compute tokens** section of the notebook.

Click **Check my progress** to verify the objective.



Task 6. Advanced features


Function calling

[Function calling](#) lets you provide a set of tools that it can use to respond to the user's prompt. You create a description of a function in your code, then pass that description to a language model in a request. The response from the model includes the name of a function that matches the description and the arguments to call it with.

For more examples of Function Calling, refer to [this lab](#).

- Run the **Function calling** section of the notebook.

Click **Check my progress** to verify the objective.



Use context caching

[Context caching](#) lets you store frequently used input tokens in a dedicated cache and reference them for subsequent requests. This eliminates the need to repeatedly pass the same set of tokens to a model.

Create a cache

- Run the **Create a cache** section of the notebook.


Use a cache

- Run the **Use a cache** section of the notebook.

Delete a cache

- Run the **Delete a cache** section of the notebook.

Click **Check my progress** to verify the objective.



Batch prediction

Different from getting online (synchronous) responses, where you are limited to one input request at a time, [batch predictions for the Gemini API in Vertex AI](#) allow you to send a large number of requests to Gemini in a single batch request. Then, the model responses asynchronously populate to your storage output location in [Cloud Storage](#) or [BigQuery](#).

Batch predictions are generally more efficient and cost-effective than online predictions when processing a large number of inputs that are not latency sensitive.

Prepare batch inputs

The input for batch requests specifies the items to send to your model for prediction.

Batch requests for Gemini accept BigQuery storage sources and Cloud Storage sources. You can learn more about the batch input formats in the [Batch text generation](#) page.

This lab uses Cloud Storage as an example. The requirements for Cloud Storage input are:

- File format: [JSON Lines \(JSONL\)](#)
- Located in `us-central1`
- Appropriate read permissions for the service account

Each request that you send to a model can include parameters that control how the model generates a response. Learn more about Gemini parameters in the [Experiment with parameter values](#) page.

This is one of the example requests in the input JSONL file `batch_requests_for_multimodal_input_2.jsonl`:

```
{"request":{"contents": [{"role": "user", "parts": [{"text": "List\nobjects in this image."}, {"file_data": {"file_uri": "gs://cloud-  
samples-data/generative-ai/image/office-desk.jpeg", "mime_type":  
"image/jpeg"}}]}], "generationConfig":{"temperature": 0.4}}}
```

- Run the **Prepare batch inputs** section of the notebook.

Prepare batch output location

When a batch prediction task completes, the output is stored in the location specified in your request.

- The location is in the form of a Cloud Storage or BigQuery URI prefix, for example: `gs://path/to/output/data` or `bq://projectId.bqDatasetId`.
- If not specified, `gs://STAGING_BUCKET/gen-ai-batch-prediction` is used for Cloud Storage source and `bq://PROJECT_ID.gen_ai_batch_prediction.predictions_TIMESTAMP` is used for BigQuery source.

This lab uses a Cloud Storage bucket as an example for the output location.

You can specify the URI of your Cloud Storage bucket in `BUCKET_URI`, or, if it is not specified, a new Cloud Storage bucket in the form of `gs://PROJECT_ID-TIMESTAMP`

is be created for you.

- Run the **Prepare batch output location** section of the notebook.

Send a batch prediction request

To make a batch prediction request, you specify a source model ID, an input source and an output location where Vertex AI stores the batch prediction results.

For more, see the [Batch prediction API](#) page. You can also check the status in the console at <https://console.cloud.google.com/vertex-ai/batch-predictions>

- Run the **Send a batch prediction request** section of the notebook.

Note: it may take up to 10 minutes for your batch prediction to complete.

Retrieve batch prediction results

When a batch prediction task is complete, the output of the prediction is stored in the location specified in your request. It is also available in

`batch_job.dest.bigquery_uri` or `batch_job.dest.gcs_uri`.

Example output:

```
{"status": "", "processed_time": "2024-11-13T14:04:28.376+00:00",
"request": {"contents": [{"parts": [{"file_data": null, "text": "List
objects in this image."}, {"file_data": {"file_uri": "gs://cloud-
samples-data/generative-ai/image/gardening-tools.jpeg", "mime_type":
"image/jpeg"}, "text": null}], "role": "user"}], "generationConfig":
{"temperature": 0.4}}, "response": {"candidates": [{"avgLogprobs":
-0.10394711927934126, "content": {"parts": [{"text": "Here's a list of
the objects in the image:\n\n* **Watering can:** A green plastic
watering can with a white rose head.\n* **Plant:** A small plant
(possibly oregano) in a terracotta pot.\n* **Terracotta pots:** Two
terracotta pots, one containing the plant and another empty, stacked on
top of each other.\n* **Gardening gloves:** A pair of striped gardening
gloves.\n* **Gardening tools:** A small trowel and a hand cultivator
(hoe). Both are green with black handles."}], "role": "model"},
"finishReason": "STOP"}], "modelVersion": "gemini-2.5-flash",
"usageMetadata": {"candidatesTokenCount": 110, "promptTokenCount": 264,
"totalTokenCount": 374}}}
```

- Run the **Retrieve batch prediction results** section of the notebook.

Click **Check my progress** to verify the objective.



Retrieve batch prediction results


Check my progress

Get text embeddings

You can get text embeddings for a snippet of text by using `embed_content` method. All models produce an output with 768 dimensions by default. However, some models give users the option to choose an output dimensionality between 1 and 768. See [Vertex AI text embeddings API](#) for details.

- Run the **Get text embeddings** section of the notebook.

Click **Check my progress** to verify the objective.



Congratulations!

Congratulations! You've successfully explored the Google Gen AI SDK, learning to connect to AI services, send diverse prompts, and fine-tune responses from the Gemini model. You've also got hands-on experience with more advanced techniques like managing interactions, using context caching, and even working with embeddings! Now you're well-equipped to leverage the power of generative AI for your own projects.

Next steps / learn more

Check out the following resources to learn more about Gemini:

- [Gemini Overview](#)
- [Generative AI on Vertex AI Documentation](#)
- [Generative AI on YouTube](#)
- Explore the Vertex AI [Cookbook](#) for a curated, searchable gallery of notebooks for Generative AI.
- Explore other notebooks and samples in the [Google Cloud Generative AI repository](#).

Google Cloud training and certification

...helps you make the most of Google Cloud technologies. [Our classes](#) include technical skills and best practices to help you get up to speed quickly and continue your learning journey. We offer fundamental to advanced level training, with on-demand, live, and virtual options to suit your busy schedule. [Certifications](#) help you validate and prove your skill and expertise in Google Cloud technologies.

Lab Last Tested August 4, 2025

Copyright 2025 Google LLC. All rights reserved. Google and the Google logo are trademarks of Google LLC. All other company and product names may be trademarks of the respective companies with which they are associated.

[< Previous](#)

[Next >](#)